


# NEURAL NETWORKS AND DEEP LEARNING



# INTRODUCTION




# ADVANTAGES OF ANNs OVER TRADITIONAL MACHINE LEARNING ALGORITHMS

The best performance

Scale effectively with data

No need for feature engineering

Adaptable and transferable





# ADVANTAGES OF TRADITIONAL MACHINE LEARNING ALGORITHMS OVER ANNS



# HIERARCHICAL DATA REPRESENTATION

- One reason that ANNs are able to perform so well is that a large number of layers allows the network to learn representations of the data at many different levels.
- At lower levels of the model, simple features are learned, such as edges and gradients, as can be seen by looking at the features that were learned in the initial layers.



# NEURAL NETWORKS AND DEEP LEARNING

This lesson covers

- Convolutional neural networks for image classification
- TensorFlow and Keras – frameworks for building neural networks
- Using pre-trained neural networks
- Internals of a convolutional neural network
- Training a model with transfer learning
- Data augmentations – the process of generating more training data



# FASHION CLASSIFICATION

The plan for our project is the following:

- First, we'll download the dataset and use a pre-trained model to classify images.
- Then we'll talk about neural networks, see how they work internally.
- After that, we'll take the pre-trained neural network and adjust it for solving our tasks.
- Finally, we'll expand our dataset by generating many many more images from the images we have.

# GPU VS CPU

- Training a neural network is a computationally demanding process, and it requires powerful hardware to make it faster.
- To speed up training, we usually use GPUs – graphical processing units, or, simply, graphic cards.
- For this lesson, a GPU is not required, You can do everything on your laptop, but without a GPU, it will be approximately 8 times slower than with a GPU.

# DOWNLOADING THE CLOTHING DATASET

- For this project, we need a dataset of clothes.
- We will use a subset of the clothing dataset[3], which contains around 3,800 images of ten different classes.
- The data is available in a GitHub repository. Let's clone it:

```
!git clone  
https://github.com/fenago/clothing-dataset-small-master.git
```

The dataset is already split into train, validation and test



Each of these folders has 10 subfolders: one subfolder for each type of clothes

Images in the dataset are organized in subfolders





0a7e5fe0-  
d592-40e6-  
b9b8-75a...



0ad5bfb5-  
0f2b-4396-  
8c05-39c...



0c2eb9ff-  
7f26-492d-  
9957-0d8...



0c99f0b4-  
3a0d-4d24-  
bfdd-e9e...



0c224954-  
0e0f-4caa-  
82c8-cf95...



0ccc318a-  
7d69-4d7f-  
a442-aac...



0db5a848-  
2066-436f-  
bd21-8b3...



0e3d71f8-  
7677-4cd4-  
ba24-478...



0e27351a-  
13d0-41a6-  
b731-409...



0e684087-  
83bf-4153-  
90f4-6f7...



0fe5eeb6-  
316f-4f60-  
b604-8a1...



1a08f33a-  
2ff4-4fb8-  
920b-8ff5...



1c9a6bc0-  
d29e-4e2b-  
8d00-d19...



1ca03195-  
b1e8-4c47-  
85de-81a...



1caba263-  
088a-448a-  
be6d-300...



1d42c614-  
19d5-4515-  
8ef9-0bf...



1d226430-  
72e6-4be7-  
998e-604...



1d407629-  
87e5-4702-  
b9fc-e3b...

# TENSORFLOW AND KERAS

- If you use your laptop with Anaconda, or run the code somewhere else, you need to install TensorFlow – a library for building neural networks.
- Use Pip to do it:

```
!pip install tensorflow
```

# TENSORFLOW AND KERAS

- We begin by importing NumPy and Matplotlib:

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

- Next, import TensorFlow and Keras:

```
import tensorflow as tf
from tensorflow import keras
```

# LOADING IMAGES

- There's a special function in Keras for loading images. It's called `load_img`.
- Let's import it:

```
from tensorflow.keras.preprocessing.image import load_img
```



# LOADING IMAGES

- Let's use this function to take a look at one of the images:

```
path =  
'/home/jovyan/clothing-dataset-small/train/t-shirt'  
name = '5f0a3fa0-6a3d-4b68-b213-72766a643de7.jpg'  
fullname = path + '/' + name  
load_img(fullname)
```

- After executing the cell, we should see an image of a T-shirt

```
path = './clothing-dataset-small/train/t-shirt'  
name = '5f0a3fa0-6a3d-4b68-b213-72766a643de7.jpg'  
fullname = path + '/' + name  
load_img(fullname)
```



- To resize the image, specify the `target_size` parameter:

```
load_img(fullname, target_size=(299, 299))
```

- As a result, the image becomes square and a bit squashed

```
load_img(fullname, target_size=(299, 299))
```



# CONVOLUTIONAL NEURAL NETWORKS

- Neural networks is a class of machine learning models for solving classification and regression problems.
- Our problem is a classification problem – we need to determine the category of an image.
- However, our problem is special: we're dealing with images.

# USING A PRE-TRAINED MODEL

- First, we'll need to import the model itself and some helpful functions:

```
from tensorflow.keras.applications.xception
import Xception
from tensorflow.keras.applications.xception
import preprocess_input
from tensorflow.keras.applications.xception
import decode_predictions
```

# USING A PRE-TRAINED MODEL

- Let's load this model:

```
model = Xception(  
    weights='imagenet',  
    input_shape=(299, 299, 3)  
)
```

# USING A PRE-TRAINED MODEL

- First, we load it using `load_img` function:

```
img = load_img(fullname, target_size=(299, 299))
```



- The `img` variable is an `Image` object, which we need to convert to a NumPy array, It's easy to do:

```
x = np.array(img)
```

- This array should have the same shape as the image, Let's check it:

```
x.shape
```

# USING A PRE-TRAINED MODEL





# USING A PRE-TRAINED MODEL

- Since we just one image, we need to create a batch with this single image:

```
X = np.array([x])
```

- : If we had several images, for example, x, y and z, we'd write:

```
X = np.array([x, y, z])
```

- Let's check its shape:

```
X.shape
```

# USING A PRE-TRAINED MODEL

- Before we can apply the model to our image, we need to prepare it.
- We do it with the `preprocess_input` function:

```
X = preprocess_input(X)
```

# GETTING PREDICTIONS

- To apply the model, use the predict method:

```
pred = model.predict(X)
```

- Let's take a look at this array:

```
pred.shape
```

# GETTING PREDICTIONS

- Luckily, there's a function `decode_predictions` that we can use to decode the prediction into meaningful class names:


```
decode_predictions(pred)
```

# GETTING PREDICTIONS

- It shows the top-5 most likely classes for this image:


```
[[('n02667093', 'abaya', 0.028757658),  
  ('n04418357', 'theater_curtain', 0.020734021),  
  ('n01930112', 'nematode', 0.015735716),  
  ('n03691459', 'loudspeaker', 0.013871926),  
  ('n03196217', 'digital_clock', 0.012909736)]]
```

# INTERNALS OF THE MODEL



# CONVOLUTIONAL LAYERS

- Even though “convolutional layer” sounds complicated, it’s nothing more than a set of filters – small “images” with simple shapes like stripes



# CONVOLUTIONAL LAYERS





No match




Feature map

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	3	3	2	2	0	0
0	0	0	0	3	3	2	2	0	0
0	0	0	0	0	0	9	9	0	0
0	0	0	0	0	0	9	9	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

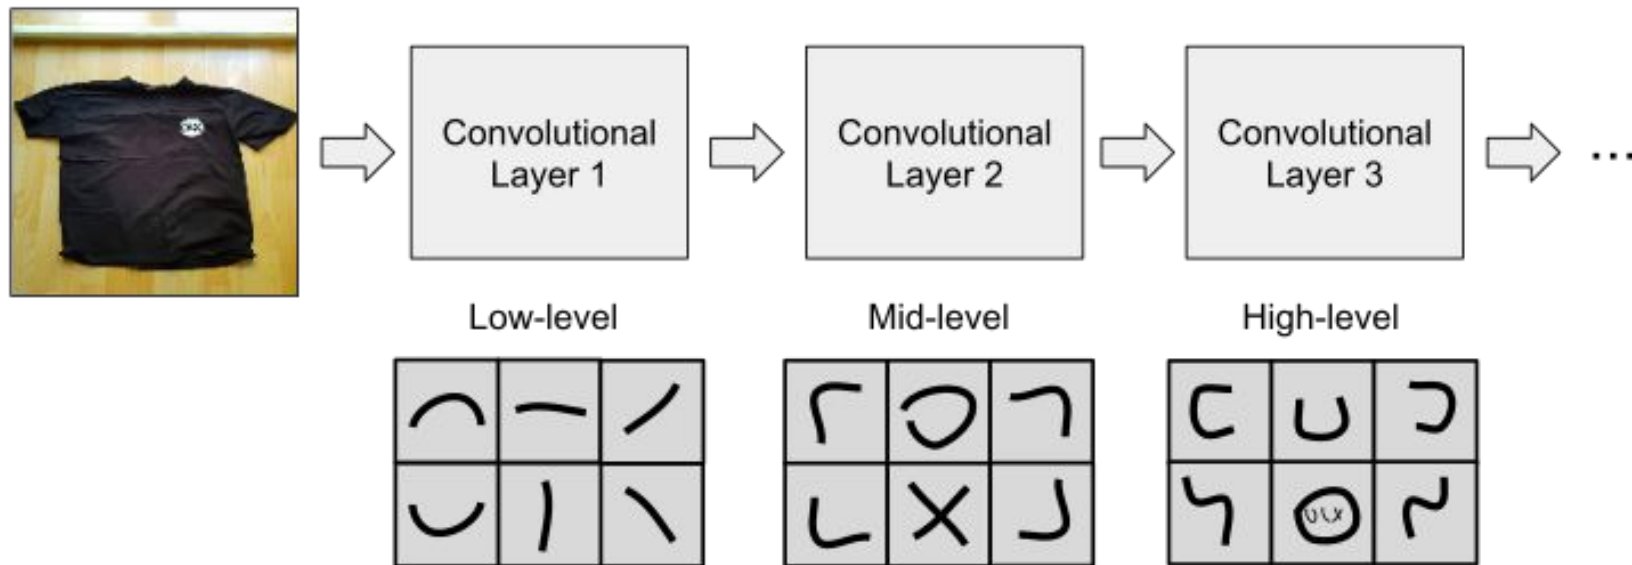


Good match

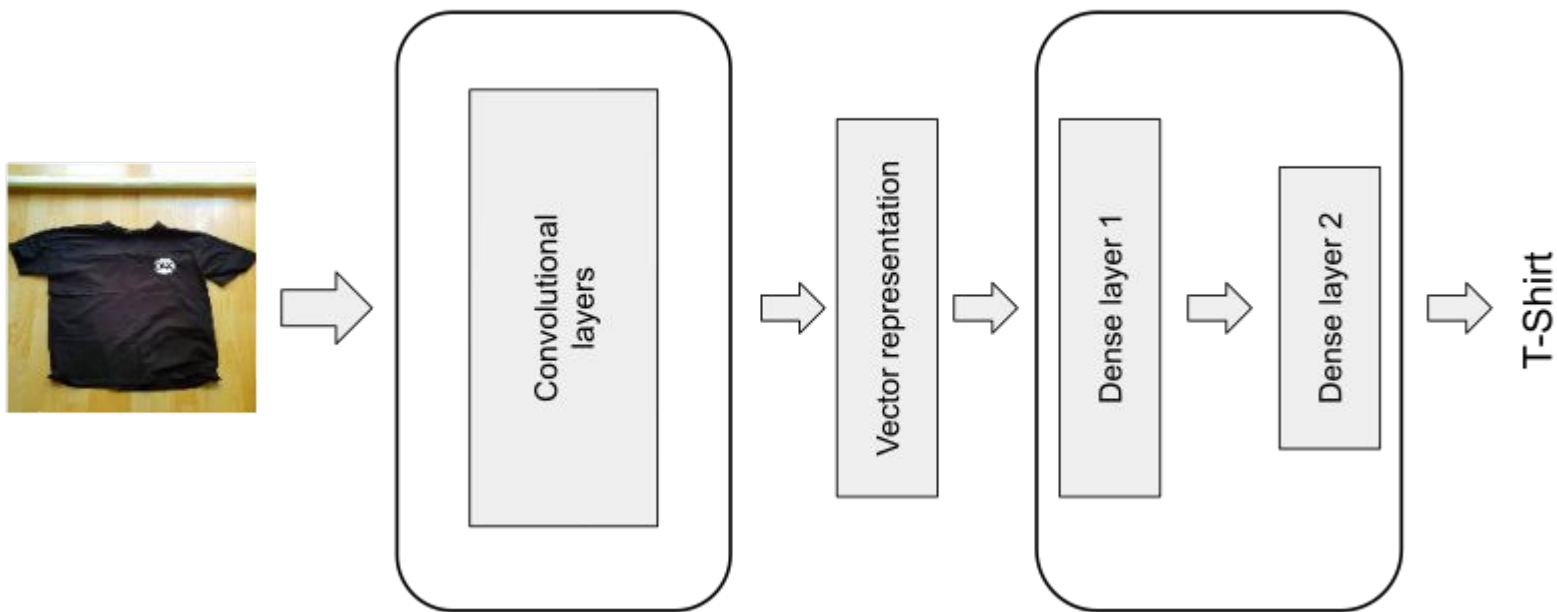
# CONVOLUTIONAL LAYERS




# CONVOLUTIONAL LAYERS




# DENSE LAYERS





# DENSE LAYERS






DENSE LAYERS







Vector representation



Dense layer 1  
(inner)

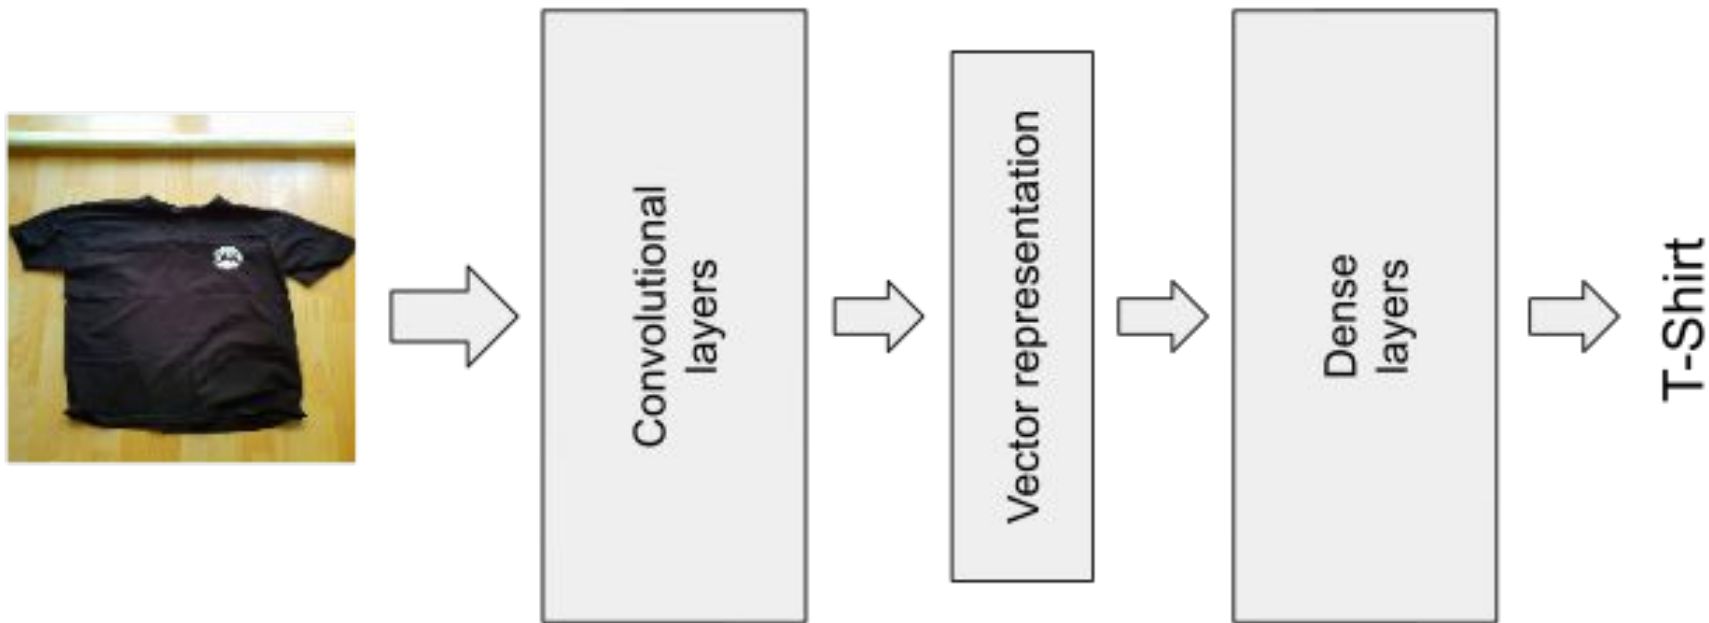


Dense layer 2  
(output)



T-Shirt


# DENSE LAYERS



# TRAINING THE MODEL

- Training a convolutional neural network takes a lot of time and requires a lot of data.
- But there's a shortcut: we can use transfer learning, an approach where we take a pre-trained model and adapt it to our problem.

# TRANSFER LEARNING



# LOADING THE DATA

- Keras comes with a solution – ImageDataGenerator.
- Instead of loading the entire dataset into memory, it loads the images from disk in small batches, Let's use it:

```
from tensorflow.keras.preprocessing.image import  
ImageDataGenerator
```

```
train_gen = ImageDataGenerator(  
    preprocessing_function=preprocess_input  
)
```

# LOADING THE DATA

- We have a generator now, so we just need to point it to the directory with the data.
- For that, use the `flow_from_directory` method:

```
train_ds = train_gen.flow_from_directory(  
    "clothing-dataset-small/train",  
    target_size=(150, 150),  
    batch_size=32,  
)
```

# LOADING THE DATA

- When we execute the cell, it informs us how many images are there in the train dataset and how many classes:

Found 3068 images belonging to 10 classes.

# LOADING THE DATA

- Now we repeat the same process for the validation dataset:

```
validation_gen = ImageDataGenerator(  
    preprocessing_function=preprocess_input  
)
```

```
val_ds = validation_gen.flow_from_directory(  
    "clothing-dataset-small/validation",  
    target_size=image_size,  
    batch_size=batch_size,  
)
```



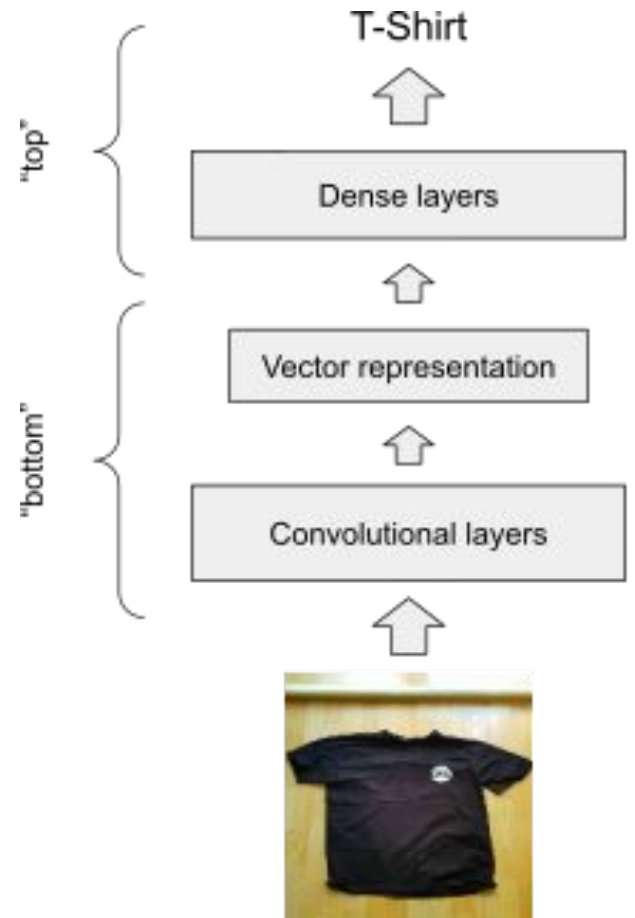
# CREATING THE MODEL

- let's create the base model:

```
base_model = Xception(  
    weights='imagenet',  
    include_top=False  
    input_shape=(150, 150, 3),  
)
```



# CREATING THE MODEL



# CREATING THE MODEL

- We don't want to train the base model: attempting to do it will destroy all the filters.
- So, we “freeze” the the base model by setting the trainable parameter to False:

```
base_model.trainable = False
```

# CREATING THE MODEL

- Now let's build the clothes classification model:

```
inputs = keras.Input(shape=(150, 150, 3))
```

```
base = base_model(inputs, training=False)
```

```
vector =
```

```
keras.layers.GlobalAveragePooling2D()(base)
```

```
outputs = keras.layers.Dense(10)(vector)
```

```
model = keras.Model(inputs, outputs)
```

# CREATING THE MODEL

- First, we specify the input and the size of the arrays we expect:

```
inputs = keras.Input(shape=(150, 150, 3))
```

- Next, we create the base model:

```
base = base_model(inputs, training=False)
```

# CREATING THE MODEL

- The result is `base`, which is a functional component (like `base_model`) that we can combine with other components.
- We use it as the input to the next layer:

```
vector =  
keras.layers.GlobalAveragePooling2D()(base)
```

# CREATING THE MODEL

- It may be a bit confusing because we create a layer and immediately connect it to base.
- We can rewrite it to make it simpler to understand:

```
pooling = keras.layers.GlobalAveragePooling2D()  
vector = pooling(base)
```

# CREATING THE MODEL

- Another functional component that we connect to the next layer – a dense layer:

```
outputs = keras.layers.Dense(10)(vector)
```



# CREATING THE MODEL

- In our case, the data comes in inputs and goes out of outputs.
- We just need to do one final step: wrap both inputs and outputs into a Model class:

```
model = keras.Model(inputs, outputs)
```

# CREATING THE MODEL



**inputs** = keras.Input(shape=(150, 150, 3))

**base** = base\_model(**inputs**, training=False)


**vector** = keras.layers.GlobalAveragePooling2D()(**base**)

**outputs** = keras.layers.Dense(10)(**vector**)

model = keras.Model(**inputs**, **outputs**)

T-Shirt

# CREATING THE MODEL



# TRAINING THE MODEL

- Let's create it:

```
learning_rate = 0.01  
optimizer = keras.optimizers.Adam(learning_rate)
```

# TRAINING THE MODEL

- We need to classify clothes into 10 different classes, so we'll use the categorical cross-entropy loss:

```
loss =  
keras.losses.CategoricalCrossentropy(from_logits=  
True)
```

# TRAINING THE MODEL

- In this case, we explicitly tell the network to output probabilities: softmax is similar to sigmoid, but for multiple classes.
- Then the output is not “logits” anymore, so we can drop this parameter:

```
loss = keras.losses.CategoricalCrossentropy()
```

# TRAINING THE MODEL

- Now let's put the optimizer and the loss together.
- For that, we'll use the compile method of our model:

```
model.compile(  
    optimizer=optimizer,  
    loss=loss,  
    metrics=["accuracy"]  
)
```

# TRAINING THE MODEL

- Our model is ready for training! To do it, use the fit method:

```
model.fit(train_ds, epochs=10,  
validation_data=val_ds)
```





- When we start training, Keras informs us about the progress:

Train for 96 steps, validate for 11 steps

Epoch 1/10

96/96 [=====] - 22s 227ms/step -  
loss: 1.2372 - accuracy: 0.6734 - val\_loss: 0.8453 -  
val\_accuracy: 0.7713

Epoch 2/10

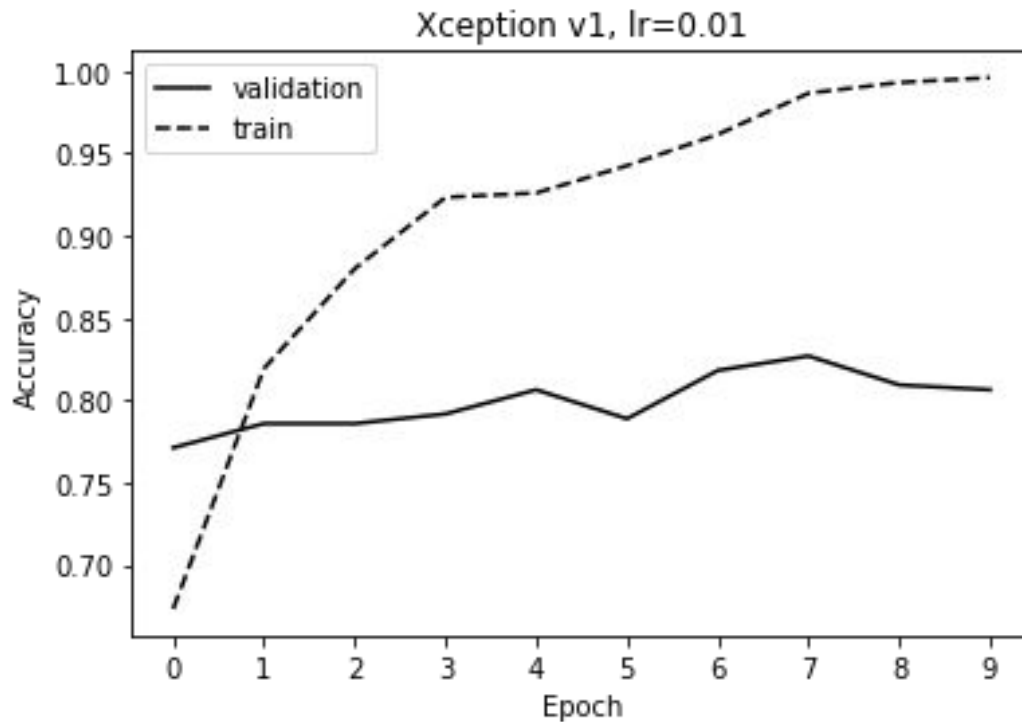
96/96 [=====] - 16s 163ms/step -  
loss: 0.6023 - accuracy: 0.8194 - val\_loss: 0.7928 -  
val\_accuracy: 0.7859

...

Epoch 10/10

96/96 [=====] - 16s 165ms/step -  
loss: 0.0274 - accuracy: 0.9961 - val\_loss: 0.9342 -  
val\_accuracy: 0.8065

# TRAINING THE MODEL



# "COMPLETE EXERCISE"



# ADJUSTING THE LEARNING RATE

```
def make_model(learning_rate):  
    base_model = Xception(  
        weights='imagenet',  
        input_shape=(150, 150, 3),  
        include_top=False  
    )  
  
    base_model.trainable = False  
  
    inputs = keras.Input(shape=(150, 150, 3))  
  
    base = base_model(inputs, training=False)  
    vector = keras.layers.GlobalAveragePooling2D()(base)  
  
    outputs = keras.layers.Dense(10)(vector)  
  
    model = keras.Model(inputs, outputs)  
  
    optimizer = keras.optimizers.Adam(learning_rate)  
    loss = keras.losses.CategoricalCrossentropy(from_logits=True)  
  
    model.compile(  
        optimizer=optimizer,  
        loss=loss,  
        metrics=["accuracy"],  
    )  
  
    return model
```

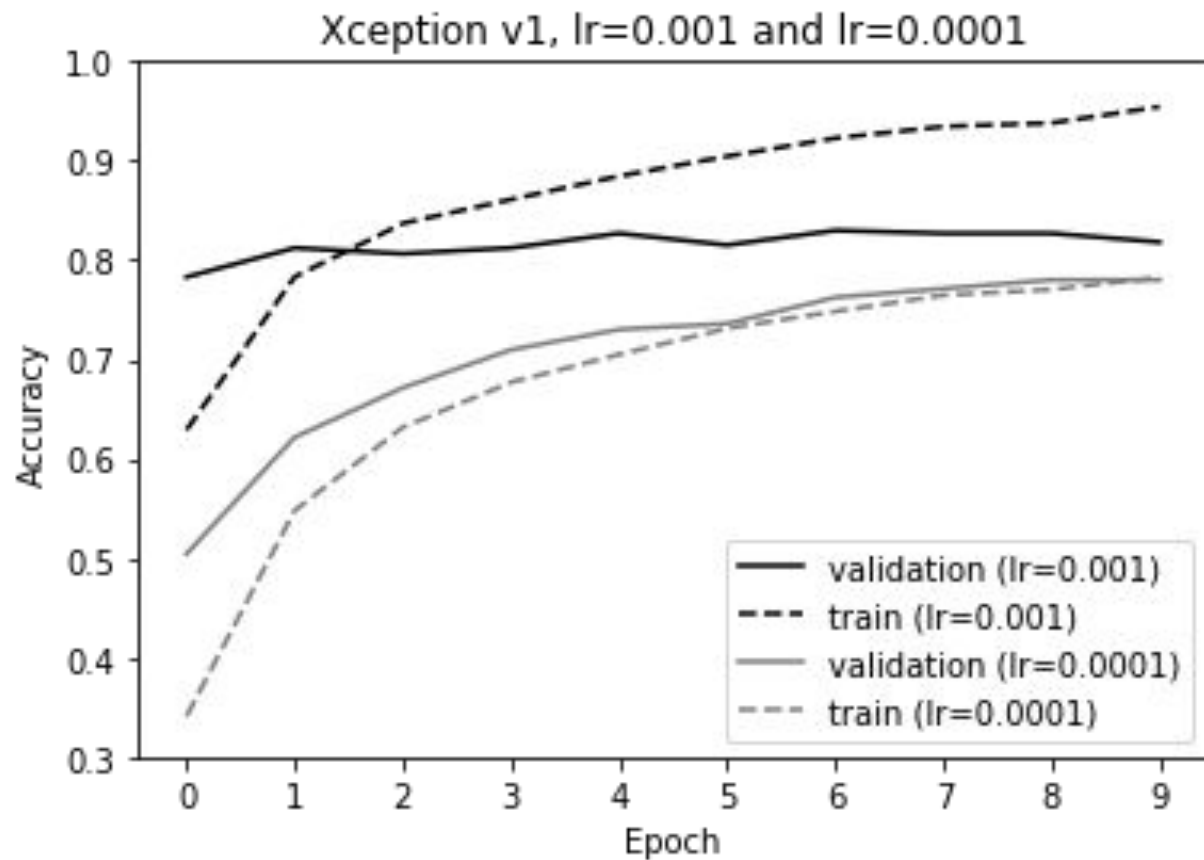
# ADJUSTING THE LEARNING RATE

- We've tried 0.01, so let's try 0.001:


```
model = make_model(learning_rate=0.001)
model.fit(train_ds, epochs=10,
validation_data=val_ds)
```

- We can also try even smaller value of 0.0001:

```
model = make_model(learning_rate=0.0001)
model.fit(train_ds, epochs=10,
validation_data=val_ds)
```



Xception v1, different learning rates



# ADJUSTING THE LEARNING RATE

- For the learning rate of 0.001, the best accuracy is 83%

Learning rate	0.01	0.001	0.0001
Validation accuracy	82.7%	83.0%	78.0%



# SAVING THE MODEL AND CHECKPOINTING

- Once the model is trained, we can save it using the `save_weights` method:

```
model.save_weights('xception_v1_model.h5',  
save_format='h5')
```

# SAVING THE MODEL AND CHECKPOINTING

- Keras has a special class for doing it:  
ModelCheckpoint, Let's use it:

```
checkpoint = keras.callbacks.ModelCheckpoint(  
    "xception_v1_{epoch:02d}_{val_accuracy:.3f}.h5",  
    save_best_only=True,  
    monitor="val_accuracy"  
)
```

- The first parameter is a template for the filename,  
Let's take a look at it again:

```
"xception_v1_{epoch:02d}_{val_accuracy:.3f}.h5"
```

# SAVING THE MODEL AND CHECKPOINTING

- We can use it by passing it to the callbacks argument of the fit method:

```
model = make_model(learning_rate=0.001)
```

```
model.fit(  
    train_ds,  
    epochs=10,  
    validation_data=val_ds,  
    callbacks=[checkpoint]  
)
```

# SAVING THE MODEL AND CHECKPOINTING



<input type="checkbox"/>	0	/	Name ↓	Last Modified	File size
<input type="checkbox"/>			clothing-dataset-small	2 days ago	
<input type="checkbox"/>			chapter-07-neural-nets.ipynb	Running seconds ago	549 kB
<input type="checkbox"/>			xception_v1_01_0.765.h5	2 minutes ago	84 MB
<input type="checkbox"/>			xception_v1_02_0.789.h5	2 minutes ago	84 MB
<input type="checkbox"/>			xception_v1_03_0.809.h5	2 minutes ago	84 MB
<input type="checkbox"/>			xception_v1_06_0.830.h5	a minute ago	84 MB

# ADDING MORE LAYERS

- we trained a model with one dense layer:

```
inputs = keras.Input(shape=(150, 150, 3))
```

```
base = base_model(inputs, training=False)
```


```
vector =
```

```
keras.layers.GlobalAveragePooling2D()(base)
```

```
outputs = keras.layers.Dense(10)(vector)
```

```
model = keras.Model(inputs, outputs)
```

```
keras.Model(inputs, outputs)
```



# ADDING MORE LAYERS

- For example, we can add a dense layer of size 100:

```
inputs = keras.Input(shape=(150, 150, 3))  
base = base_model(inputs, training=False)  
vector = keras.layers.GlobalAveragePooling2D()(base)
```

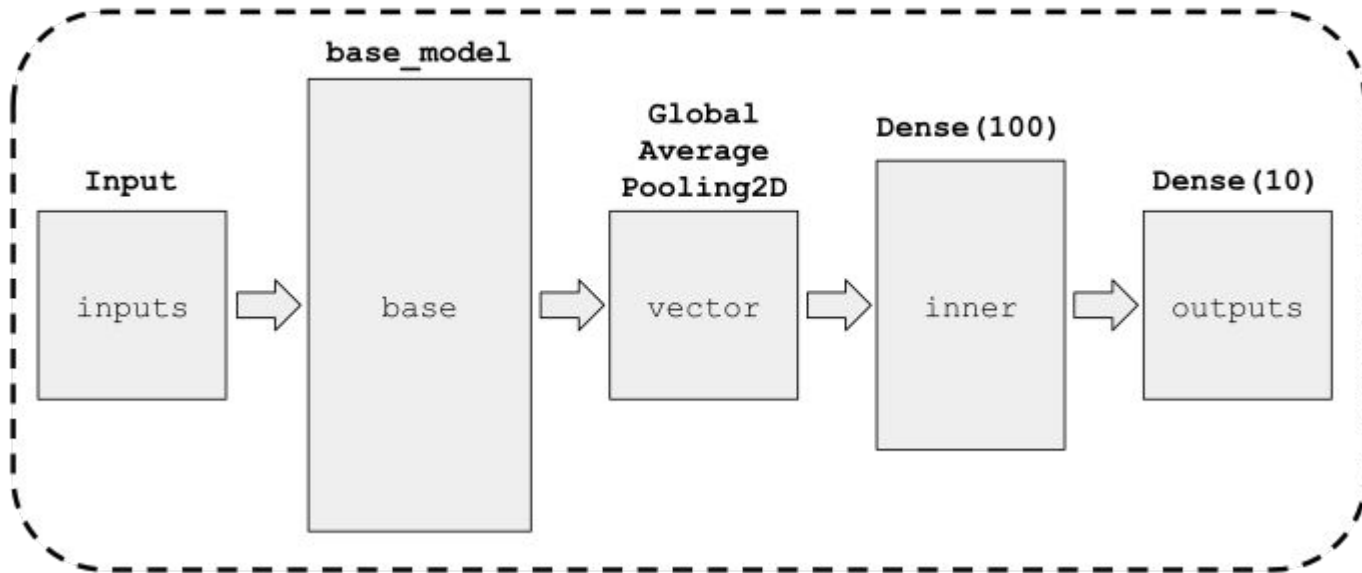
```
inner = keras.layers.Dense(100,  
activation='relu')(vector)
```

```
outputs = keras.layers.Dense(10)(inner)
```

```
model = keras.Model(inputs, outputs)
```

# ADDING MORE LAYERS

```
keras.Model(inputs, outputs)
```






# ADDING MORE LAYERS


- Let's take another look at the line with the new dense layer:

```
inner = keras.layers.Dense(100,  
activation='relu')(vector)
```

- Here, we set the activation parameter to “relu”.



(A) Two dense layers  
without dropout



(B) Two dense layers  
with dropout

# REGULARIZATION AND DROPOUT

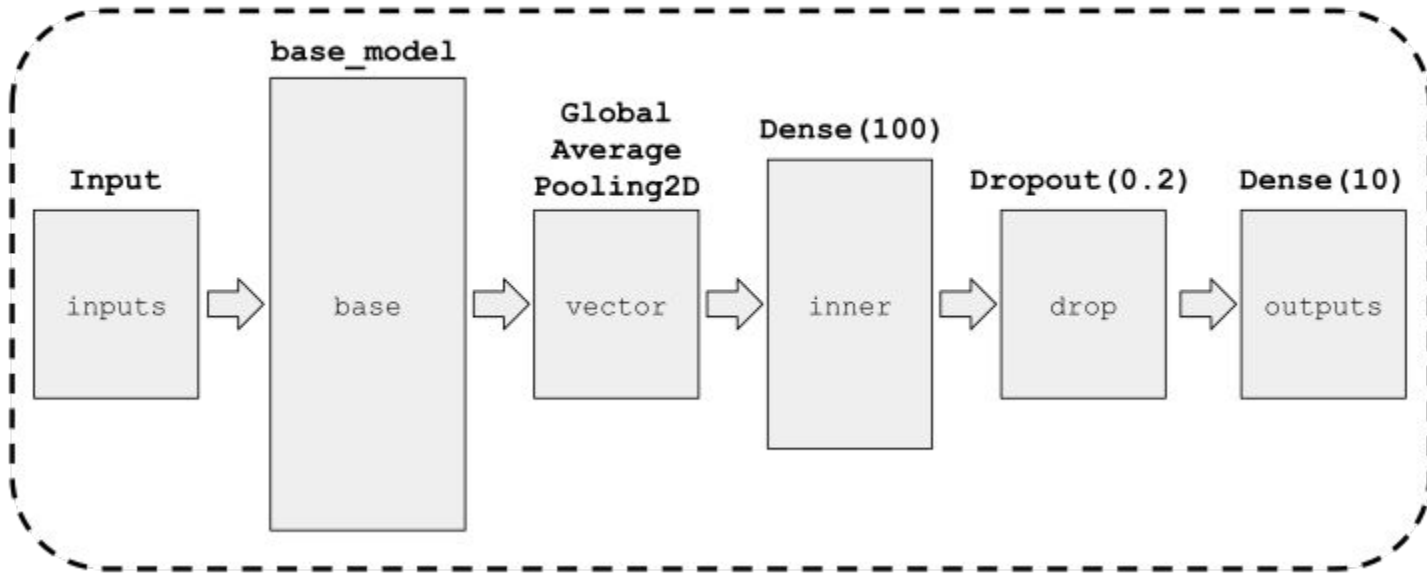
```
inputs = keras.Input(shape=(150, 150, 3))  
base = base_model(inputs, training=False)  
vector =  
keras.layers.GlobalAveragePooling2D()(base)
```

```
inner = keras.layers.Dense(100,  
activation='relu')(vector)  
drop = keras.layers.Dropout(0.2)(inner)  
outputs = keras.layers.Dense(10)(drop)
```

```
model = keras.Model(inputs, outputs)
```

# REGULARIZATION AND DROPOUT

```
keras.Model(inputs, outputs)
```



```

def make_model(learning_rate, droprate):
    base_model = Xception(
        weights='imagenet',
        input_shape=(150, 150, 3),
        include_top=False
    )

    base_model.trainable = False

    inputs = keras.Input(shape=(150, 150, 3))
    base = base_model(inputs, training=False)
    vector = keras.layers.GlobalAveragePooling2D()(base)

    inner = keras.layers.Dense(100, activation='relu')(vector)
    drop = keras.layers.Dropout(droprate)(inner)

    outputs = keras.layers.Dense(10)(drop)

    model = keras.Model(inputs, outputs)

    optimizer = keras.optimizers.Adam(learning_rate)
    loss = keras.losses.CategoricalCrossentropy(from_logits=True)

    model.compile(
        optimizer=optimizer,
        loss=loss,
        metrics=["accuracy"],
    )

    return model


```

# REGULARIZATION AND DROPOUT

- So, let's train it:

```
model = make_model(learning_rate=0.001,  
droprate=0.0)  
model.fit(train_ds, epochs=30,  
validation_data=val_ds)
```

Xception v2, different dropout rates




# REGULARIZATION AND DROPOUT

- The best accuracy we could achieve is 84.5% for the dropout rate of 0.5

Dropout rate	0.0	0.2	0.5	0.8
Validation accuracy	84.2%	84.2%	84.5%	82.4%



Xception v2, different dropout rates (train)





"COMPLETE  
EXERCISE "

# DATA AUGMENTATION

- The process of generating more data from an existing dataset is called data augmentation



# DATA AUGMENTATION

- The easiest way to create a new image from an existing one is to flip it horizontally, vertically, or both



# DATA AUGMENTATION

- Rotating is another image manipulation strategy that we can use: we can generate a new image by rotating an existing one by some degree

rotation=-30



rotation=-15



rotation=0



rotation=15



rotation=30



# DATA AUGMENTATION

- When the shear is positive, we pull the right side down, and when it's negative, we pull the right side up

shear=-20



shear=-10



shear=0



shear=10



shear=20



Shear



shear -30




shear 30




Rotation



rotate -30



rotate 30



Shifting an image horizontally. Positive values shift the image to the left, while negative values shift it to the right

width\_shift=-30



width\_shift=-15



width\_shift=0



width\_shift=15



width\_shift=30



Shifting an image vertically. Positive values shift the image to the top, while negative values shift it to the bottom

height\_shift=-30



height\_shift=-15



height\_shift=0



height\_shift=15



height\_shift=30





# DATA AUGMENTATION

- Finally, we can zoom an image in or out

zoom=0.5



zoom=0.75



zoom=1



zoom=1.2



zoom=1.5



# DATA AUGMENTATION



# DATA AUGMENTATION

- For example, we can create a new generator:

```
train_gen = ImageDataGenerator(  
    rotation_range=30,  
    width_shift_range=30.0,  
    height_shift_range=30.0,  
    shear_range=10.0,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    vertical_flip=False,  
    preprocessing_function=preprocess_input  
)
```

# DATA AUGMENTATION

- For our project, we'll take a small set of these augmentations:

```
train_gen = ImageDataGenerator(  
    shear_range=10.0,  
    zoom_range=0.1,  
    horizontal_flip=True,  
    preprocessing_function=preprocess_input,  
)
```

# DATA AUGMENTATION

- We use the generator in the same way as previously:

```
train_ds = train_gen.flow_from_directory(  
    "clothing-dataset-small/train",  
    target_size=(150, 150),  
    batch_size=32,  
)
```

# DATA AUGMENTATION

- We load the validation dataset using exactly the same code as before:

```
validation_gen = ImageDataGenerator(  
    preprocessing_function=preprocess_input  
)
```

```
val_ds = validation_gen.flow_from_directory(  
    "clothing-dataset-small/validation",  
    target_size=image_size,  
    batch_size=batch_size,  
)
```

# DATA AUGMENTATION

- We're ready to train a new model now:

```
model = make_model(learning_rate=0.001,  
droprate=0.2)  
model.fit(train_ds, epochs=50,  
validation_data=val_ds)
```

# "COMPLETE EXERCISE"





# TRAINING A LARGER MODEL

- Now we're ready to train a model!

```
model = make_model(learning_rate=0.001,  
droprate=0.2)  
model.fit(train_ds, epochs=20,  
validation_data=val_ds)
```

# USING THE MODEL (LOADING THE MODEL)

 alexeygrigorev released this 13 hours ago · [3 commits](#) to master since this release

Pre-trained models for chapter 7 - detecting types of clothes

## ▼ Assets 4

 <a href="#">xception_v3_44_0.853.h5</a>	82.2 MB
 <a href="#">xception_v4_large_08_0.894.h5</a>	82.2 MB
 <a href="#">Source code (zip)</a>	
 <a href="#">Source code (tar.gz)</a>	

# LOADING THE MODEL

- To use it, load the model using the `load_model` function from the `models` package:

```
model =  
keras.models.load_model('xception_v4_large_08_0.8  
94.h5')
```

# EVALUATING THE MODEL

- To load the test data we follow the same approach: we use ImageDataGenerator, but point to the “test” directory, Let’s do it:

```
test_gen = ImageDataGenerator(  
    preprocessing_function=preprocess_input  
)
```

```
test_ds = test_gen.flow_from_directory(  
    "clothing-dataset-small/test",  
    shuffle=False,  
    target_size=(299, 299),  
    batch_size=32,  
)
```

# EVALUATING THE MODEL

- Evaluating a model in Keras is as simple as invoking the evaluate method:

```
model.evaluate(test_ds)
```

- It applies the model to all the data in the test folder and shows the evaluation metrics: loss and accuracy.

```
12/12 [=====] - 70s  
6s/step - loss: 0.2493 - accuracy: 0.9032
```

# EVALUATING THE MODEL

- If we repeat the same process for the small dataset, we'll see that the performance is worse:

12/12 [=====] - 15s  
1s/step - loss: 0.6931 - accuracy: 0.8199

# GETTING THE PREDICTIONS

- If we want to apply the model to a single image, we need to do the same thing ImageDataGenerator perform internally:
  1. load an image
  2. pre-process it
- We already know how to load an image. We can use `load_img` for that:

```
path =  
'clothing-dataset-small/test/pants/c8d21106-bbdb-4e8d-8  
3e4-bf3d14e54c16.jpg'  
img = load_img(path, target_size=(299, 299))
```

```
img = load_img(path, target_size=(299, 299))  
img
```





# GETTING THE PREDICTIONS

- Next, we pre-process the image:

```
x = np.array(img)
X = np.array([x])
X = preprocess_input(X)
```

- And, finally, get the predictions:

```
pred = model.predict(X)
```

# GETTING THE PREDICTIONS

- We can see the predictions for the image by checking the first row of predictions: `pred[0]`

```
pred = model.predict(X)
pred[0]
```

```
array([-2.8609195, -4.234049 , -1.573255 , -1.9078847, 10.24705 ,  
       -2.2489128, -4.297381 ,  4.43905  , -4.458805 , -3.9616926],  
      dtype=float32)
```

# GETTING THE PREDICTIONS

- To get the element with the highest score, we can use the `argmax` method.
- It returns the index of the element with the highest score

```
pred[0].argmax()
```

```
4
```

```
labels = {  
    0: 'dress',  
    1: 'hat',  
    2: 'longsleeve',  
    3: 'outwear',  
    4: 'pants',  
    5: 'shirt',  
    6: 'shoes',  
    7: 'shorts',  
    8: 'skirt',  
    9: 't-shirt'  
}
```

## GETTING THE PREDICTIONS

# GETTING THE PREDICTIONS

- To get the label, simply look it up in the dictionary:

```
labels[pred[0].argmax()]
```





# "COMPLETE EXERCISES & LAB"

# SUMMARY



- TensorFlow is a framework for building and using neural networks. Keras is a library on top of TensorFlow that makes training models simpler.
- For image processing, we need a special kind of neural networks: convolutional neural networks.
- They consist of a series of convolutional layers followed by a series of dense layers.